

Does language change the outcome? Tables and figures (generated from results/main-1 and results/predict-1, commit 420cdbc; regenerate with scripts/paper_figures.py and scripts/results_package.py).

Table 1. Main run: what each model decided in Greek, its agreement with the court, and the pre-specified rates. F, W and L are the noise floor, rewording rate and language rate, each a share of the ten cases comparing 31 draws with 31 (Greek halves of 62); $\Delta_EN = L_EN - W_en$, paired bootstrap over cases (10,000 resamples). Mean D is the grant share in English minus the grant share across all Greek draws, averaged over cases. "Ground differs" is the share of comparisons, among those where the modal outcome held, in which the modal ground differed: Greek against English, and Greek against the back-translation.

Model	Dismissed (Greek)	Granted, full or part (Greek)	Agreement with court	MCC	F	W_en	L_EN	Δ_EN [95% CI]	W_de	L_DE	Mean D [95% CI]	Ground differs, EN / BT-en
GPT-5.6 Terra	100%	0%	0.40	0.00	0.00	0.00	0.00	0.00 [0.00, 0.00]	0.00	0.00	0.00 [0.00, 0.01]	0.10 / 0.00
Terra, reasoning off	96%	4%	0.40	0.00	0.00	0.00	0.00	0.00 [0.00, 0.00]	0.00	0.00	-0.04 [-0.11, 0.00]	0.40 / 0.20
Claude Sonnet 5	88%	12%	0.40	0.00	0.00	0.20	0.00	-0.20 [-0.50, 0.00]	0.10	0.00	-0.10 [-0.18, -0.04]	0.36 / 0.11
Gemini 3.6 Flash	93%	7%	0.50	0.29	0.00	0.00	0.00	0.00 [0.00, 0.00]	0.10	0.10	-0.01 [-0.03, 0.01]	0.00 / 0.00

Table 2. Decide against predict. The main run (the model decides as a judge; 31 draws per arm, 62 Greek) and the exploratory predictive run (the model predicts the court's decision as a legal analyst; 15 draws per arm, 30 Greek), Greek and English arms. Flips are cases whose Greek and English modal outcomes differ. The predictive run had no back-translation arm, so its language rate is not net of rewording.

Model	Dismissed (Greek), decide	Agreement with court, decide	F, decide	L_EN, decide	Mean D [95% CI], decide	Flips to refusing in EN, decide	Dismissed (Greek), predict	Agreement with court, predict	F, predict	L_EN, predict	Mean D [95% CI], predict	Flips to refusing in EN, predict
GPT-5.6 Terra	100%	0.40	0.00	0.00	0.00 [0.00, 0.01]	0	100%	0.40	0.00	0.00	0.02 [0.00, 0.06]	0
Terra, reasoning off	96%	0.40	0.00	0.00	-0.04 [-0.11, 0.00]	0	100%	0.40	0.00	0.00	0.00 [0.00, 0.00]	0
Claude Sonnet 5	88%	0.40	0.00	0.00	-0.10 [-0.18, -0.04]	0	65%	0.50	0.18	0.33	-0.31 [-0.47, -0.15]	3
Gemini 3.6 Flash	93%	0.50	0.00	0.00	-0.01 [-0.03, 0.01]	0	92%	0.50	0.01	0.01	-0.00 [-0.03, 0.03]	0

Table 3. Considerations invoked in the models' reasoning, as the share of all valid replies in the main run (10 cases $\times$ 6 arms per model) in which an automated keyword screen, in Greek, English and German, matched the consideration. A screen, not hand coding.

Consideration invoked in the reasoning (share of valid replies)	GPT-5.6 Terra	Terra, reasoning off	Claude Sonnet 5	Gemini 3.6 Flash
disputed facts / conflicting evidence	74%	73%	82%	71%
credibility / cross-examination	11%	11%	63%	3%
trial or full hearing needed	34%	46%	77%	57%
fanciful / realistic distinction	19%	18%	16%	16%
documents / expert evidence needed	7%	9%	11%	4%
real prospect (the limb (a) phrase)	76%	75%	70%	99%
compelling reason (the limb (b) phrase)	18%	27%	22%	6%
mini-trial	0%	0%	0%	0%



Figure 2. Grant share per case (share of draws granting the application in full or in part) under the decide framing (filled markers) and the predict framing (hollow), Greek (blue) and English (orange), one panel per model. Shaded columns are the cases the court granted in full or in part.